

# Traveling Salesman Problem (TSP)



# Team Members & Roles

- **Eng: Shahd Mohamed Hamed Abdelrahman**

- **Task:** Implementation of **Uniform-Cost Search (UCS)**.
- **Responsibilities:** Developing an uninformed search strategy that guarantees optimal solutions by expanding nodes based on the lowest cumulative path cost.

- **Eng: Omar Mohamed Ibrahim Badawy**

- **Task:** Implementation of **A\* Search Algorithm**
- **Responsibilities:** Designing an informed search using path cost plus a heuristic estimate ( $f(n) = g(n) + h(n)$ ) to reduce search space and improve speed.

- **Eng: Noor Hussain Mwafi**

- **Task:** Implementation of **Hill Climbing Search**
- **Responsibilities:** Developing a local optimization framework using **Random-restart** and **Swap-based neighbor generation** to iteratively improve tour quality.

- **Eng: Nour Yasser Hashem El-Sheikh**

- **Task:** Implementation of **Nearest Neighbor + 2-opt**.
- **Responsibilities:** Combining a greedy constructive algorithm with a local search operator (2-opt) to remove "crossover" edges and achieve very fast, near-optimal solutions

- **Eng: Hend Mohamed Mohamed Fiala**

- **Task:** Implementation of **Genetic Algorithm (GA)**
- **Responsibilities:** Applying evolutionary concepts (Selection, Crossover, and Mutation) to maintain a population of solutions and perform a global search of the solution space

# INTRODUCTION

## Navigating the Traveling Salesman Problem (TSP)

The **Traveling Salesman Problem (TSP)** stands as one of the most iconic challenges in combinatorial optimization and computer science. The objective is deceptively simple: find the shortest possible route that visits a set of cities exactly once and returns to the origin. However, due to its **NP-hard** nature, the search space expands factorially, making brute-force solutions impossible for large-scale maps.

To tackle this complexity, our team has implemented a diverse suite of algorithmic strategies—ranging from exact uninformed searches to sophisticated meta-heuristics. Each approach offers a unique trade-off between computational speed and solution optimality.

### 1. Systematic Search Foundations

- **Uniform-Cost Search (UCS):** Led by **Eng. Shahd Mohamed**, this implementation focuses on an uninformed search strategy. By expanding nodes based on the lowest cumulative path cost, it establishes a rigorous baseline that guarantees an optimal solution.
- *A Search Algorithm:*\* Developed by **Eng. Omar Mohamed**, this informed search enhances efficiency by integrating path cost with a heuristic estimate ( $f(n) = g(n) + h(n)$ ), significantly pruning the search space without sacrificing accuracy.

### 2. Local Optimization & Heuristics

- **Hill Climbing Search:** **Eng. Noor Hussain** implements a local optimization framework. By utilizing **Random-restart** and **Swap-based** neighbor generation, this approach iteratively climbs toward better solutions, effectively navigating the landscape of local optima.
- **Nearest Neighbor + 2-opt:** **Eng. Nour Yasser** focuses on a hybrid approach. This combines a greedy constructive algorithm for immediate results with a **2-opt** local search operator to eliminate "crossover" edges, achieving high-quality solutions with remarkable speed.

### 3. Evolutionary Intelligence

- **Genetic Algorithm (GA):** Managed by **Eng. Hend Mohamed**, this strategy applies the principles of natural selection. Through **Selection, Crossover, and Mutation**, the algorithm maintains a diverse population of potential routes, performing a global search to find near-optimal tours in vast solution spaces.

## 1. Uniform Cost Search (UCS)

The Uniform-Cost Search (UCS) is an uninformed search strategy designed to find the optimal path in a weighted graph<sup>1</sup>.

- **Core Strategy:** It expands nodes based on the lowest cumulative path cost from the source.
- **Optimality:** Unlike greedy heuristics, UCS guarantees the discovery of the shortest possible route for any given TSP instance.
- **Objective:** Its primary role in this project is to provide a gold-standard "optimal" baseline to measure the quality of other heuristic methods.

---

### Phase 1: Exploration and Expansion

In this phase, the algorithm explores the search space systematically without any prior knowledge of city locations.

- **Mechanism:** It maintains a priority queue of paths, always choosing the one with the least total distance to expand next.
- **Exhaustive Nature:** It explores all possible directions equally until the goal state (visiting all cities and returning) is reached
- **Uniformity:** Because it does not use a heuristic, it treats all neighboring cities solely based on their actual distance from the current node.

---

### Phase 2: Guaranteeing Optimality

The strength of UCS lies in its mathematical certainty regarding the final result.

- **Optimal Solution:** It is proven to find the global minimum path cost.
- **Comparison to Heuristics:** While methods like Nearest Neighbor or Hill Climbing might settle for a "good" path, UCS continues until it confirms no shorter path exists.
- **Cost Focus:** Every decision is based on the actual path cost ( $g(n)$ ), ensuring the most efficient route is never overlooked<sup>10</sup>.

## 1. Uniform Cost Search (UCS)

### Best Performance Conditions:

- **Absolute Precision:** Best used when the **optimal solution** (shortest path) is strictly required.
- **Small-Scale Tasks:** Highly effective for small datasets where the state space is manageable.
- **Optimal Baseline:** Ideal for creating a reference point to evaluate the quality of other heuristic algorithms.

### Factors Affecting Performance:

- **Computational Expense:** Performance drops drastically as the number of cities increases due to **extremely slow** execution.
- **Memory Exhaustion:** It suffers from **extremely high memory usage** because it stores all explored paths.
- **Lack of Guidance:** Unlike A\*, it operates without heuristics, exploring the search space in all directions, which limits its scalability.

Number of Cities (N)	Execution Time (Speed)	Memory Usage	Solution Quality (Optimality)
N = 5 (Small)	Fast	Manageable	100% Optimal (Best Choice)
N = 15 (Medium)	Slow	Very High	100% Optimal
N = 20 (Large/Complex)	Extremely Slow	Critical (May crash)	100% Optimal (Impractical)

## 2.A\* Search

A\* is an Informed Search Algorithm used to find the shortest path between nodes. It combines the strengths of Uniform Cost Search (BFS with weights) and Greedy Best-First Search.

### ● The Core Formula

The algorithm evaluates nodes using the fitness function:

$$f(n) = g(n) + h(n)$$

- **$g(n)$ :** The actual cost from the start node to the current node n.
- **$h(n)$ :** The **Heuristic** estimated cost from node n to the goal (the "informed" part).
- **$f(n)$ :** Total estimated cost of the cheapest solution through node n.

### Mechanism

A\* uses a **Priority Queue** (Min-Heap) to always expand the node with the lowest  $f(n)$  value. This ensures that the algorithm explores the most promising paths first, significantly reducing the number of nodes visited compared to Uninformed Search.

---

### ● What makes A\* Perfect (Optimal)? (Conditions)

To guarantee the shortest path, the Heuristic  $h(n)$  must satisfy two conditions:

1. **Admissibility:** The heuristic must **never overestimate** the actual cost to reach the goal. It must be optimistic or exact ( $h(n) < h^*(n)$ ).
2. **Consistency (Monotonicity):** For every node n and its successor  $n'$ , the estimated cost to the goal from n is no greater than the step cost to  $n'$  plus the estimated cost from  $n'$ .

### ● Factors that Reduce Efficiency (Conditions)

1. **Poor Heuristic:** If  $h(n)$  is too low (close to 0), A\* behaves like BFS (very slow, explores everything).
  2. **State Space Explosion:** In problems like TSP, as the number of cities increases, the number of possible paths grows factorially ( $n!$ ), leading to high **Memory Consumption**.
  3. **Inadmissible Heuristic:** If  $h(n)$  overestimates, the algorithm might find a solution faster, but it **won't be the optimal (shortest)** solution.
-

## 2. A\* Search

In the Traveling Salesperson Problem (TSP), A\* struggles with "Complexity" as the number of cities grows. Here is a comparison of what happens during the evaluation:

Cities 20	Cities 15	Cities 5	Metric
<b>Slow</b> :Performance drops significantly (Exponential growth).	<b>Moderate</b> :Can take seconds to a few minutes depending on CPU.	<b>Instant</b> :Results in milliseconds.	<b>Execution Time</b>
<b>Critical</b> :Likely to hit memory limits or cause "Out of Memory" errors.	<b>High</b> :Storing unvisited states consumes significant RAM.	<b>Very Low</b> :The priority queue stays small.	<b>Memory Usage</b>
<b>Absolute Minimum</b> : If the search completes.	<b>Absolute Minimum</b> : Guaranteed to be the shortest possible route.	<b>Absolute Minimum</b> :Finds the perfect shortest path.	<b>Path Cost (Tour Length)</b>
<b>Imbalance</b> :Speed becomes too slow to maintain optimality.	<b>Shift</b> :Maintaining optimality starts to sacrifice speed.	<b>Perfect Balance</b> : High speed and 100% optimal.	<b>Optimality vs. Speed</b>
<b>Poor</b> :Not scalable for larger real-world datasets.	<b>Struggling</b> : Reaching the limits of exact search methods.	<b>Excellent</b> :Handles small sets with ease.	<b>Scalability</b>

The provided Python code implements A\* for TSP using:

1. **heapq**: To manage the Priority Queue based on the lowest  $f(n)$ .
2. **distance function**: Calculates the Euclidean distance ( $g(n)$ ).
3. **heuristic function**: Uses the distance to the nearest unvisited city as a simple estimation to guide the search.
4. **solve method**: Continuously pops the best path, checks if all cities are visited, and returns to the start city to complete the cycle.

### 3. Hill Climbing Algorithm

#### Algorithm Mechanics, Challenges, and Implementation

- **Local Optimum & Random Restarts:**

- **The Trap:** A **Local Optimum** occurs when the algorithm reaches a state better than its immediate neighbors but inferior to the best possible global solution.
- **The Solution:** **Random Restarts** act as a safeguard; by re-launching the search from different random starting points, we increase the probability of finding the **Global Optimum**.
- **Scaling:** As the number of cities increases, the landscape becomes more "rugged" with more local traps, necessitating an increase in the number of restarts.

- **Reproducibility (Seed vs. Restart):**

- **Random Seed:** Ensures the "randomness" is deterministic. It allows different users to get the **exact same results** every time they run the code, which is essential for scientific evaluation.
- **Random Restart:** Operates *within* a single run to explore various regions of the search space. The Seed simply ensures that these "random" regions are identical across different executions.

- **Technical Features:**

- **Heuristic Calculation:** Real-world distances are calculated in **Kilometers (km)** using the **Haversine Formula**.
- **Visualization:** Interactive **Folium** maps highlight the **Start City in Green** and provide city names and visitation order upon clicking.

### 3. Hill Climbing Algorithm

#### Performance Metrics, Conditions for Success, and Conclusion

- **Comparative Analysis (5 vs. 15 vs. 20 Cities):**

- **Execution Time:** Increases with city count and restarts; 5 cities are near-instant, while 20 cities require more computation time to ensure quality.
- **Memory Usage:** Extremely efficient ( $O(n)$ ) across all cases as it only stores the current and best paths.
- **Path Cost & Optimality:** 5 cities consistently find the global optimum. For 15-20 cities, the algorithm provides a high-quality approximation (Near-Optimal) rather than a guaranteed perfect solution.

- **Optimal Conditions for Hill Climbing:**

- **Best Scenario:** Small to medium datasets ( $N < 30$ ) where speed is prioritized over absolute perfection.
- **Landscape:** Performs best when the search space has a clear gradient (slope) toward the shorter path

Impact of Increasing N	Cities 20	Cities 15	Cities 5	Metric
<b>Very Fast</b> .Even with more cities, it only checks a few neighbor states.	0.01~s	0.005~s	s0.0005~	<b>Execution Time</b>
<b>Extremely Low</b> .It only stores the "Current Path" and the "Best Path".	Minimal.	Minimal.	.Minimal	<b>Memory Usage</b>
<b>Quality decreases</b> .With 20 cities, there are too many "traps" (local optima).	Sub-optimal	Fair / Variable.	.Optimal	<b>Path Cost</b>
Great for speed, but reliability drops as the search space grows.	Low Accuracy.	Good.	.Excellent	<b>Optimality vs Speed</b>

## 4. Nearest Neighbor + 2-opt

### Overview of the Hybrid Approach

This project utilizes a hybrid heuristic strategy that combines construction and optimization.

- **Methodology:** It integrates the Nearest Neighbor (NN) algorithm with 2-opt Local Search.
  - **Primary Objectives:** The goal is to achieve fast execution, low memory consumption, and a near-optimal solution.
  - **Practicality:** This approach offers a balanced trade-off between the quality of the solution and computational efficiency.
- 

### Phase 1: Construction via Nearest Neighbor (NN)

The process begins by generating a feasible starting tour.

- **Mechanism:** Starting from a chosen city, the algorithm moves to the closest unvisited city until the cycle is complete.
  - **Advantages:** It is simple to implement, extremely fast, and requires very little memory.
  - **Role:** It serves as a baseline for further improvement, though it is greedy and does not guarantee an optimal solution.
- 

### Phase 2: Refinement via 2-opt Local Search

The rough initial tour is refined to remove obvious inefficiencies.

- **Mechanism:** The algorithm selects two edges and reverses the segment between them to see if it reduces the total distance.
- **Edge Crossing Removal:** A key purpose of this stage is to eliminate "edge crossings" in the tour.
- **Termination:** The process repeats iteratively until no further local improvements can be found

## 4. Nearest Neighbor + 2-opt

### Best Performance Conditions:

- **Scalability:** Best suited for moderate to large-scale TSP instances where fast solutions are required.
- **Efficiency:** Ideal when memory consumption must be kept to a minimum.
- **Goal:** Preferred when a high-quality, near-optimal solution is acceptable rather than a strictly optimal one.

### Factors Affecting Performance:

- **Starting Point:** The quality of the final tour is highly dependent on the initial city selected during the NN phase.
- **Local Optimum Trapping:** The algorithm stops once no nearby improvements exist, meaning it can miss the global optimal solution.
- **Problem Size:** For very small problems, its performance is "worse" than UCS or A\* because it doesn't guarantee absolute optimality.

Number of Cities (N)	Execution Time (Speed)	Memory Usage	Solution Quality (Optimality)
N = 5 (Small)	Instantaneous	Minimal	Near-Optimal (UCS/A* are better for absolute optimality)
N = 15 (Medium)	Very Fast	Low	High Quality (Best balance of speed and precision)
N = 20 (Large/Complex)	Highly Efficient	Consistently Low	Reliable (May hit local optima, but stays practical)

## 5. Genetic Algorithm (GA)

### Overview of Evolutionary Intelligence

The Genetic Algorithm (GA) represents a meta-heuristic approach inspired by the process of natural selection.

- **Global Search:** Unlike local search methods, GA explores the entire solution space to find near-optimal routes.
- **Population-Based:** It maintains a diverse set of potential solutions (tours) simultaneously.
- **Goal:** To achieve high-quality results in large-scale TSP problems where other methods might fail.

---

### Phase 1: Population and Selection

The process begins by creating and filtering potential solutions.

- **Initialization:** A population of random or heuristic-based tours is generated to start the evolution.
- **Selection:** Solutions with shorter path lengths (higher fitness) are prioritized for "reproduction."
- **Role:** This phase ensures that the best traits of current routes are passed down to future generations.

---

### Phase 2: Crossover and Mutation

This stage introduces variation and exploration to improve the results.

- **Crossover:** Parts of two parent tours are combined to create "offspring," attempting to inherit the best segments of both.
- **Mutation:** Small, random changes are made to individual tours to maintain diversity and prevent the algorithm from getting stuck in local optima.
- **Purpose:** This simulates biological evolution to explore new regions of the search space.

## 5. Genetic Algorithm (GA)

### Evaluation and Best Performance

GA is evaluated based on its ability to handle complex problems.

- **Best Performance:** It is most effective for large-scale TSP instances where global exploration is required.
- **Resource Cost:** It requires more execution time and computational resources compared to simpler heuristics like NN.
- **Optimality:** While it doesn't guarantee the absolute mathematical optimum like UCS, it is much more likely to find a global optimum than Hill Climbing.

Performance Comparison Table

Number of Cities (N)	Execution Time	Memory Usage	Solution Quality
Small (N=5)	Slow (Compared to NN/A*)	Moderate	High (Overkill for small sets)
Medium (N=15)	Moderate	High	Very High (Very reliable results)
Large (N=20+)	Best Choice	Intensive	Global Optimum (Best at avoiding local traps)

### GA Performance Conditions

- **Best Performance Conditions:** Excels in large-scale problems <sup>8</sup>and scenarios where global exploration is necessary to avoid local optima<sup>9</sup>.
- **Conditions Affecting Performance:** The size of the population, the mutation rate, and the available computational time directly impact the quality of the final solution

# Performance Metrics Overview

Metric	UCS	A*	Hill Climbing	NN + 2-Opt	Genetic (GA)
Memory Usage	Extremely High	High	Minimal	Minimal	Moderate
Scalability	Very Poor	Limited	Excellent	Excellent	Excellent
Best Case	Small Data	Precision Needs	Quick Drafts	Industrial Use	Large-scale TSP
Path Cost (Length)	Near-Minimum	Minimum (Optimal)	Minimum (Optimal)	Near-Minimum	Minimum (Optimal)

## Best Algorithm For This Problem

From Algorithms We Used

NN + 2-Opt (Why?)

Best practical choice because it provides a near-optimal solution almost instantly by effectively removing path intersections.

At All

LKH (Lin-Kernighan) (Why ?)

Best overall globally because it is the world-standard algorithm capable of solving massive TSP problems with thousands of cities at extreme speed.

# Comparative Analysis by Number of Cities

## 5 Cities (Small)

- **All algorithms** perform well here.
- **UCS & A\***: Find the absolute shortest path instantly (only  $4! = 24$  permutations).
- **Complexity**: Negligible for all.

## 15 Cities (Medium)

- **UCS**: Starts to struggle. The search tree expands to millions of nodes, consuming significant RAM.
- **A\***: Still optimal and efficient **if** using a strong heuristic like MST (Minimum Spanning Tree).
- **NN + 2-Opt**: Finds a solution almost identical to A\* but in a fraction of a second.
- **Genetic**: Requires several generations to converge, might be slower than 2-Opt for this size.

## 20 Cities (The Breaking Point)

- **UCS**: Likely to **crash** or run out of memory.  $19!$  is a massive number of states.
- **A\***: Becomes very slow; memory consumption becomes the primary bottleneck.
- **Hill Climbing**: Often gets "stuck" in a sub-optimal path (Local Optimum) and fails to find the best route.
- **NN + 2-Opt**: This is the **sweet spot**. It provides a near-optimal solution (within 1-3% of the best) almost instantly.
- **Genetic**: Highly effective here. It explores diverse areas of the map and usually beats Hill Climbing in quality.

---

---

**Thank You For Helping Me**

